

RW_GNN Code: Functional Overview and Workflow

Yunxaing Wang

June 22, 2025

1 Overall Workflow and Structure

Algorithm 1 Quick start with RW_GNN

```
python rwgnn_dis_s6.py --dataset Cora \
    --alpha 5.0 --c 0.01 \
    --num-walks-per-node 10 \
    --epochs 50 --walk-bluetype levy
```

1. **Argument Parsing:** Use `argparse` to read dataset name, learning rate, dropout, weight decay, random-walk hyperparameters (walks per node, random starts per graph), early stopping patience, etc.
2. **Environment Setup:** Fix NumPy and PyTorch seeds, select CPU/GPU device, create checkpoint directory.
3. **Data Loading:**
 - If `--dataset Cora`, instantiate `LocalCoraDataset`: read raw `ind.cora.*` files, merge and re-order features/labels, build `edge_index`, split into 140/500/1000 train/val/test masks with no overlap or empty sets.
 - Otherwise, use `torch_geometric.datasets.TUDataset`, and apply one-hot degree encoding for graphs without node attributes.
4. **Model Construction:** Instantiate `RW_GNN`, which encapsulates `ExplicitRandomWalkEncoder`, passing all hyperparameters; set up Adam optimizer and a `ReduceLROnPlateau` scheduler.
5. **Training and Evaluation:**
 - **Cora:** Single-graph training over epochs, record train/val/test loss and accuracy, save best model, trigger early stopping, and emit overfitting warnings when validation loss rises.
 - **TUDataset:** 10-fold cross-validation, 90% train / 10% validation per fold, save best model each fold, report average test accuracy across folds.
6. **Visualization:** Call `plot_training_curves` to generate and save a PNG of loss and accuracy curves side by side.

2 Key Classes and Functions

2.1 Class `ExplicitRandomWalkEncoder`

`__init__(input_dim, hidden_dim, walk_length, num_walks_per_start, device, num_random_starts_per_graph)`

Initialize the encoder: set up a two-layer feature encoder with ReLU and Dropout, a two-layer GRU with dropout, a distance cache, and all hyperparameters.

`clear_cache()` Clear the cached distance matrices when parameters like `alpha` or `walk_type` change.

compute_graph_distances(edge_index, num_nodes, graph_nodes) Compute the shortest-path distance matrix and distance-based transition probabilities for a subgraph, with caching.

- **edge_index**: edge list of the subgraph $[2 \times E]$.
- **graph_nodes**: global node indices of the subgraph $[N]$.
- Returns (walk_probs, distance_matrix).

_compute_simple_distance_matrix(adj, num_nodes, max_distance) Approximate distances via 1/2/3-hop adjacency powers.

_compute_shortest_path_distances_scipy(adj, num_nodes, max_distance) Use SciPy Dijkstra for exact distances up to a cutoff; fallback to approximation on failure.

_compute_shortest_path_distances_networkx(adj, num_nodes, max_distance) Use NetworkX all-pairs for small graphs, sampling-based approximation for large graphs.

_compute_approximate_distances_networkx(G, num_nodes, max_distance) Approximate large-graph distances by sampling single-source shortest paths and filling gaps.

_compute_walk_probabilities(distance_matrix, alpha) Build a normalized transition matrix from distances, mix with uniform jumps via damping factor c .

sample_random_walks_for_graph_classification(edge_index, batch, num_nodes) For each graph in the batch, sample multiple random walks per start node, returning walk sequences and their graph IDs.

forward(x, edge_index, batch) Execute sampling, feature encoding, GRU encoding; return {path_encodings, walk_batch} for downstream aggregation.

2.2 Class RW_GNN

__init__(input_dim, max_step, hidden_graphs, size_hidden_graphs, hidden_dim, penultimate_dim, normalization) Initialize the GNN: wrap the random-walk encoder and build a two-layer downstream network with ReLU, BatchNorm, and Dropout.

init_weights() Uniformly initialize hidden-graph parameters (reserved for future extensions).

forward(data) Dispatch to node or graph classification based on **data.batch**.

forward_node_classification(data) Node-level forward: compute per-node walk encodings, apply $fc1 \rightarrow bn \rightarrow dropout \rightarrow fc2 \rightarrow dropout2$, output log-softmax.

_node_specific_walk_encoding(x, edge_index, num_nodes) For each node, sample multiple biased or simple walks, encode via GRU, and average to obtain a node embedding.

_generate_biased_walk(start_node, walk_probs, walk_length) Generate one random walk using the provided transition probabilities.

_generate_single_walk(start_node, adj_list, walk_length) Generate one pure random walk by uniform neighbor sampling.

_simple_neighborhood_encoding(x, edge_index, num_nodes) Fallback one/two-hop feature aggregation when walk-based encoding fails.

forward_graph_classification(data) Graph-level forward: pool path encodings by graph ID, apply downstream network, output log-softmax.

2.3 Auxiliary Functions

plot_training_curves(train_losses, val_losses, train_accs, val_accs, save_path) Plot and save side-by-side training vs. validation loss and accuracy curves.

train(model, loader, optimizer, device) Perform one epoch of training, return average loss and accuracy.

test(model, loader, device) Perform one epoch of evaluation (no gradient), return average loss and accuracy.

2.4 Class LocalCoraDataset

`raw_file_names / processed_file_names` Specify raw and processed file lists (PyG interface).

`download()` No-op to disable auto-download.

`process()` 1. Read `.x`, `.y`, `.allx`, `.ally`, `.graph`, `.test.index`.

2. Merge and reorder features/labels, build an undirected `edge_index`.

3. Split into 140/500/1000 train/val/test masks, replenishing randomly if needed; enforce no overlap or empty sets.

4. Save the resulting `Data` object to disk.

`_read_pickle_file / _read_test_index` Load raw data from Pickle files or test index text.